A Translation Floor for Cross-Lingual Archives: Spelling and Content Analysis in *Imvo Zabantsundu*

Gopolang Mohlakola 2680747
University of the Witwatersrand, 1 Jan Smuts Ave,
Braamfontein, Johannesburg, 2000, South Africa

Tawananyasha Kadango 2588206
University of the Witwatersrand, 1 Jan Smuts Ave,
Braamfontein, Johannesburg, 2000, South Africa

Lily de Melo 2545080
University of the Witwatersrand, 1 Jan Smuts Ave,
Braamfontein, Johannesburg, 2000, South Africa

Abstract

In its 1884–1894 issues, *Imvo Zabantsundu* published an isiXhosa section and a smaller English one. This study asks whether they differ in content once the language signal is removed from LaBSE embeddings, and whether modernising isiXhosa spelling meaningfully changes the answer. A classifier's accuracy on the newspaper's translations, the floor, estimates how well it tells the sections apart without much difference in content. Performance above this floor suggests that the sections differ in content, mainly in subject matter. Modernising the spelling brings translations slightly closer, but its effect on the content difference is inconclusive.

Keywords: historical text normalisation; multilingual sentence embeddings; isiXhosa

Introduction

Imvo Zabantsundu, founded by John Tengo Jabavu in 1884, published mainly in isiXhosa, with some English translations and other items (Schoots, Ngwendu, et al., 2023). NLP tools are mostly developed from modern text in high-resource languages (Joshi, Santy, et al., 2020), which gives researchers a reason to modernise the spelling of isiXhosa texts first, for example changing *Tixo* and *umteto* to *Thixo* and *umthetho*. Historical normalisation is judged by spelling accuracy or task scores (Bollmann, 2019), rarely by whether it changes a conclusion, making it a "forking path" (Gelman & Loken, 2013).

Research question. *After removing the language signal from a multilingual encoder's embeddings, can a classifier still distinguish between Imvo's two sections based on their content, and does modernising the spelling change the size of this difference?* **H1:** A classifier predicting which of the two sections a chunk (a 60-word piece of a page) comes from should achieve a higher balanced accuracy (BA; chance is 0.5) on all test chunks than on translated pairs. Pairs are chunks matched to their likely translation from the other section. Each pair should share the same content, so BA on these pairs is the floor, an estimate of how well the classifier tells the sections apart without a content difference. The difference between the two BAs, named Δ, estimates the content difference. H1 is supported if the 95% confidence interval (CI) of Δ, resampled over test issues, lies above zero. **H2:** Modernising isiXhosa spelling should lower the floor, since the encoder reads modern text better, while leaving Δ unchanged within $\pm$0.02 (measurement invariance). Beyond this range, modernising changes the conclusion.

Background

Historical isiXhosa. The archive's machine-read (OCR) text was corrected by native isiXhosa speakers and deliberately not spell-checked, as 1880s spelling differs from today's (Schoots, Ngwendu, et al., 2023). Spelling is not neutral. Written isiXhosa was first shaped by missionary grammars made for learners rather than speakers (Maseko, 2017), and low-resourcedness is also a question of who builds tools and for whom (Nekoto, Marivate, et al., 2020).

Removing language. LaBSE (Feng, Yang, et al., 2022) was trained on translations in 109 languages, including isiXhosa, but its embeddings still identify the language. This signal can be removed by subtracting each language's average embedding (mean-centring; Libovický, Rosa, et al., 2020), which also removes average content differences, or by projecting out a few directions (Yang, Yang, et al., 2021; Ravfogel, Elazar, et al., 2020). As with control tasks (Hewitt & Liang, 2019), the newspaper's own translations are the control.

Methods

Data. From the 517 Imvo issues in the archive files, 511 with at least 1,000 words per language are used. Chunks in the wrong language and text repeated across issues are left out, leaving 45,465 isiXhosa and 22,927 English chunks. Issues are split 357/77/77 (train/validation/test), stratified by year.

Spelling modernisation. A word list is built from 3.3M tokens of modern isiXhosa (mostly OPUS-100 (Zhang, Williams, et al., 2020), and also Wikipedia and Vukuzenzele (Lastrucci, Rajab, et al., 2023)). *Aspiration* restores the *h* that 1880s spelling left out after letters *t, p, k, r, c, q, x*, but only when the old spelling is almost absent from the modern word list ($\leq$2 uses) and the new one is common in it ($\geq$5). *Prefix* joins prefixes written apart (*e Palamente*$\rightarrow$*ePalamente*) only when the joined form is in the word list ($\geq$5 uses). The rules rewrite 9.1% of the 3.7M isiXhosa words, including 191 likely names (capitalised in $\geq$90% of uses). *Modern* spelling applies both rules, while *raw* is the original.

Embeddings and translated pairs. Each chunk is embedded with LaBSE and, for comparison, multilingual-E5 (Wang et al., 2024). Within each issue, an isiXhosa and an English chunk that are each other's closest match by cosine similarity create a pair if their similarity is $\geq$0.55 (622 of the 10,289 test chunks form 311 test pairs). Of 100 sampled test pairs, a fluent isiXhosa speaker labelled 24 herself and reviewed 27 that a team member had labelled with AI translations and suggested labels. Of these 51, 31% were translations, 29% partial and 39% different, suggesting that the 0.55 threshold is too loose. Her own labels find more pairs different (54% vs 26%), so the reviewed ones may lean towards the AI, though Δ on pairs labelled translation or partial is similar (0.24, 0.26).

Removing the language signal (pair removal). The differences within training pairs, which are mostly language, are stacked and their top k directions are projected out of every embedding:

$$D = \left[x_i^{\text{xh}} - x_i^{\text{en}}\right]_{i \in P}, \qquad \tilde{x} = x - V_k^{\top} V_k x, \tag{1}$$

where V_k holds the top-k right singular vectors of D. A logistic regression, weighted so both sections count equally, is trained on training issues to predict each chunk's section. Two settings are chosen on validation issues: the number of directions removed, named k, is the smallest value in $\{0, 1, 2, 4, \ldots, 64\}$ that brings the floor to 0.55 or below (64 if none does, as here), and the regularisation strength $C \in \{10^{-2}, \ldots, 10^{3}\}$ is chosen for the highest BA. **Spelling effects.** For H2, each spelling

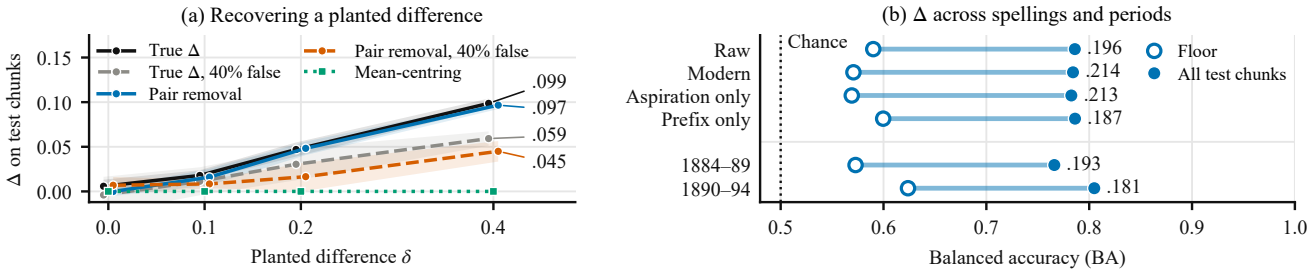


Figure 1: (a) Simulation: Δ found as the planted difference δ grows (5 seeds, ±std). Dashed lines have 40% false pairs. (b) Test issues: each line runs from the floor to BA on all test chunks, and its length is Δ.

variant keeps the pairs and k chosen from the original text, so only spelling is changed. The variant is equivalent to raw if the 90% CI of its change in Δ lies within ±0.02.

Topic clustering. Topics are k-means clusters of the language-removed embeddings, fitted without section labels. To test how much of Δ is subject matter alone, a *topic-only rule* is applied. It assigns a chunk to the isiXhosa section if its topic's isiXhosa share in training issues is above the overall share (as in Figure S1). This, and the check of where errors fall, were performed after the main results, so they are exploratory.

Synthetic check (planted gap). Each simulated chunk is a topic vector plus a language vector. A content difference is planted: English chunks use all topics evenly, while a share δ of isiXhosa chunks use only a quarter of topics. Simulated pairs copy isiXhosa content into English, and a classifier on topic vectors alone measures the true Δ. This checks the language removal and the floor, not LaBSE.

Results

The method recovers planted gaps. Tested first on simulated data with a known answer (Figure 1a), the method recovers Δ (0.097 vs true 0.099 at $\delta{=}0.4$). With 40% false pairs (unrelated chunks recorded as translations), it shrinks (0.045 vs a true 0.059). But real false pairs may behave differently, as they are paired for looking similar.

The classifier is above the floor (H1). On Imvo, the classifier reaches 0.786 on all test chunks but only 0.590 on translated pairs, so Δ is 0.196 (0.167–0.224), supporting H1. Of 311 test pairs, it gives both halves one label for 201 pairs (it cannot tell which is the translation), correctly tells them apart for 83 and labels both the wrong way round for 27. Δ is positive in every Table S2 setting (smallest 0.046) and both time periods (Figure 1b), but its size depends on the pairs used (0.117 to 0.286 across the pair sets in Table S2). No k reached the 0.55 target. Removing one direction lowered the test floor from 0.977 to 0.593, and removing more barely changed it. Removing language with only pairs of similarity ≥0.65 raised the floor to 0.725, as only 362 training pairs remain.

Topic accounts for much of Δ. A topic-only rule reaches 0.661 on test chunks, but only 0.523 on translated pairs, whose halves should share a topic. Topic accounts for about 70% of Δ (0.138 of 0.196; Table S2). Chunks on topics the other section usually covers are correctly predicted 59% of the time, against 88% for chunks on their own section's usual topics. Topic is not everything, though, as within a topic the classifier still averages 0.73 (over 24 topics).

Spelling's effect on Δ is unresolved (H2). Aspiration cuts words missing from modern text from 31.5% to 24.5%, and modernising moves translated pairs closer (similarity +0.005). Yet its effect on Δ cannot be told apart from zero (+0.018, 90% CI −0.003 to +0.037; Figure 1b, Table S1; likewise for E5 encoder), and the floor's fall (0.019) is within noise. Every CI crosses ±0.02, so H2 is neither confirmed nor rejected.

Discussion

Subject matter is the clearest difference. The sections are known to differ (Schoots, Ngwendu, et al., 2023), so what is new is measuring how much once language is removed, and showing that cleaning the spelling does not detectably change the answer. The English section leans towards colonial policy: commissions, land laws and liquor licensing; the isiXhosa section towards community life: association meetings, sport, marriage, land sales and market prices (Figure S1). The classifier errs where one section printed what the other usually did: 23 of 24 English land-sale notices go to the isiXhosa section, and most isiXhosa land-law and liquor articles go to the English one.

Translation is an imperfect control. The floor assumes that a translation says the same thing, yet the isiXhosa speaker noted the saying *isiXhosa asitolikwa*, "isiXhosa cannot be translated": even good translations adapt, turning "a feather in its cap" into the *indwe*, the crane feather of honour (Y. Siko, pers. comm.). If LaBSE misses meaning that may sometimes be attached to culture, the floor on translations may control content only roughly.

Modernisation and the historical record. Modernisation rules rewrite how people named themselves in 1884: a people (*Xosa*), a community (*Tembu*), a town (*Mtata*) and God (*Tixo*). The tools also read English more easily (2.11 vs 1.26 LaBSE tokens per word), which may tempt researchers to study only the English half (Nekoto, Marivate, et al., 2020). **Limitations.** Only 51 pairs were hand-checked and no one checked the rewrites. Topic checks use LaBSE only, and the method measures what was printed but cannot say why. Future archive studies should report the floor under both spellings.

Acknowledgements

The authors thank Yonela Siko, a fluent isiXhosa speaker, for checking the 51 translated pairs used to verify the floor.

Contribution Statements

We, Gopolang Mohlakola, Tawananyasha Kadango and Lily de Melo, worked together fairly on this project.

References

Bollmann, M. (2019). A large-scale comparison of historical text normalization systems. In *Proceedings of NAACL-HLT.*

Feng, F., Yang, Y., Cer, D., Arivazhagan, N., & Wang, W. (2022). Language-agnostic BERT sentence embedding. In *Proceedings of ACL.*

Gelman, A., & Loken, E. (2013). *The garden of forking paths* (Tech. Rep.). Columbia University.

Hewitt, J., & Liang, P. (2019). Designing and interpreting probes with control tasks. In *Proceedings of EMNLP-IJCNLP* (pp. 2733–2743).

Joshi, P., Santy, S., Budhiraja, A., Bali, K., & Choudhury, M. (2020). The state and fate of linguistic diversity and inclusion in the NLP world. In *Proceedings of the 58th annual meeting of the association for computational linguistics (ACL)* (pp. 6282–6293).

Lastrucci, R., Rajab, J., Shingange, M., Njini, D., & Marivate, V. (2023). Preparing the Vuk'uzenzele and ZA-gov-multilingual South African multilingual corpora. In *Proceedings of the fourth workshop on resources for african indigenous languages (RAIL 2023)* (pp. 18–25).

Libovický, J., Rosa, R., & Fraser, A. (2020). On the language neutrality of pre-trained multilingual representations. In *Findings of EMNLP.*

Maseko, P. (2017). Exploring the history of the writing of isiXhosa: An organic or an engineered process? *International Journal of African Renaissance Studies*, *12*(2), 81–96. doi: 10.1080/18186874.2017.1400218

Nekoto, W., Marivate, V., Matsila, T., Fasubaa, T., Fagbohungbe, T., et al. (2020). Participatory research for low-resourced machine translation: A case study in African languages. In *Findings of the association for computational linguistics: EMNLP 2020* (pp. 2144–2160). doi: 10.18653/v1/2020.findings-emnlp.195

Ravfogel, S., Elazar, Y., Gonen, H., Twiton, M., & Goldberg, Y. (2020). Null it out: Guarding protected attributes by iterative nullspace projection. In *Proceedings of ACL* (pp. 7237–7256).

Schoots, J., Ngwendu, A., De Wet, J., & Muftic, S. (2023). IsiXhosa intellectual traditions digital archive: Digitizing isiXhosa texts from 1870–1914. In *Proceedings of the fourth workshop on resources for african indigenous languages (RAIL 2023)* (pp. 54–64).

Wang, L., et al. (2024). *Multilingual E5 text embeddings: A technical report.* (arXiv:2402.05672)

Yang, Z., Yang, Y., Cer, D., & Darve, E. (2021). A simple and effective method to eliminate the self language bias in multilingual representations. In *Proceedings of EMNLP* (pp. 5825–5832).

Zhang, B., Williams, P., Titov, I., & Sennrich, R. (2020). Improving massively multilingual neural machine translation and zero-shot translation. In *Proceedings of ACL.*

NeurIPS Paper Checklist

1. **Claims**

   Question: Do the main claims made in the abstract and introduction accurately reflect the paper's contributions and scope?

   Answer: [Yes]

   Justification: The abstract and Introduction state the research question, hypotheses H1 and H2 and what we measure; Results and Discussion report and qualify each, and the Discussion ends with the study's limits, including that it measures content, not editorial intent.

   Guidelines:

   - The answer NA means that the abstract and introduction do not include the claims made in the paper.
   - The abstract and/or introduction should clearly state the claims made, including the contributions made in the paper and important assumptions and limitations. A No or NA answer to this question will not be perceived well by the reviewers.
   - The claims made should match theoretical and experimental results, and reflect how much the results can be expected to generalize to other settings.
   - It is fine to include aspirational goals as motivation as long as it is clear that these goals are not attained by the paper.

2. **Limitations**

   Question: Does the paper discuss the limitations of the work performed by the authors?

   Answer: [Yes]

   Justification: Methods, Results, Discussion: a floor that depends on the pairs used, a loose pair threshold (of the 51 hand-checked pairs, 31% were translations, 29% partial and 39% different), a simulation that does not model pairs selected for similarity, no number of removed directions reaching the 0.55 validation target, only 51 hand-checked pairs (27 of them labelled with AI suggestions and reviewed by the isiXhosa speaker, with the possible lean towards the suggestions reported), no hand-checking of normaliser rewrites, translations as an imperfect control (*isiXhosa asitolikwa*), topic checks with LaBSE only, and that the method measures what was printed but not why.

   Guidelines:

   - The answer NA means that the paper has no limitation while the answer No means that the paper has limitations, but those are not discussed in the paper.
   - The authors are encouraged to create a separate "Limitations" section in their paper.
   - The paper should point out any strong assumptions and how robust the results are to violations of these assumptions (e.g., independence assumptions, noiseless settings, model well-specification, asymptotic approximations only holding locally). The authors should reflect on how these assumptions might be violated in practice and what the implications would be.
   - The authors should reflect on the scope of the claims made, e.g., if the approach was only tested on a few datasets or with a few runs. In general, empirical results often depend on implicit assumptions, which should be articulated.
   - The authors should reflect on the factors that influence the performance of the approach. For example, a facial recognition algorithm may perform poorly when image resolution is low or images are taken in low lighting. Or a speech-to-text system might not be used reliably to provide closed captions for online lectures because it fails to handle technical jargon.
   - The authors should discuss the computational efficiency of the proposed algorithms and how they scale with dataset size.
   - If applicable, the authors should discuss possible limitations of their approach to address problems of privacy and fairness.
   - While the authors might fear that complete honesty about limitations might be used by reviewers as grounds for rejection, a worse outcome might be that reviewers discover limitations that aren't acknowledged in the paper. The authors should use their best judgment and recognize that individual actions in favor of transparency play an important role in developing norms that preserve the integrity of the community. Reviewers will be specifically instructed to not penalize honesty concerning limitations.

3. **Theory Assumptions and Proofs**

   Question: For each theoretical result, does the paper provide the full set of assumptions and a complete (and correct) proof?

   Answer: [NA]

Justification: The paper has no theoretical results.

Guidelines:

- The answer NA means that the paper does not include theoretical results.
- All the theorems, formulas, and proofs in the paper should be numbered and cross-referenced.
- All assumptions should be clearly stated or referenced in the statement of any theorems.
- The proofs can either appear in the main paper or the supplemental material, but if they appear in the supplemental material, the authors are encouraged to provide a short proof sketch to provide intuition.
- Inversely, any informal proof provided in the core of the paper should be complemented by formal proofs provided in appendix or supplemental material.
- Theorems and Lemmas that the proof relies upon should be properly referenced.

4. **Experimental Result Reproducibility**

   Question: Does the paper fully disclose all the information needed to reproduce the main experimental results of the paper to the extent that it affects the main claims and/or conclusions of the paper (regardless of whether the code and data are provided or not)?

   Answer: [Yes]

   Justification: The Introduction and Methods give the chunk size, issue filter, issue-level splits, normaliser rules and thresholds, pair selection, the language-removal equation, the hyperparameter grids and selection rules, and the resampling over test issues. The exact rule for repeated text and the E5 pair threshold are in README.txt, and the submitted code implements everything with fixed seeds. Secondary results can shift slightly with package versions, as README.txt explains.

   Guidelines:

   - The answer NA means that the paper does not include experiments.
   - If the paper includes experiments, a No answer to this question will not be perceived well by the reviewers: Making the paper reproducible is important, regardless of whether the code and data are provided or not.
   - If the contribution is a dataset and/or model, the authors should describe the steps taken to make their results reproducible or verifiable.
   - Depending on the contribution, reproducibility can be accomplished in various ways. For example, if the contribution is a novel architecture, describing the architecture fully might suffice, or if the contribution is a specific model and empirical evaluation, it may be necessary to either make it possible for others to replicate the model with the same dataset, or provide access to the model. In general. releasing code and data is often one good way to accomplish this, but reproducibility can also be provided via detailed instructions for how to replicate the results, access to a hosted model (e.g., in the case of a large language model), releasing of a model checkpoint, or other means that are appropriate to the research performed.
   - While NeurIPS does not require releasing code, the conference does require all submissions to provide some reasonable avenue for reproducibility, which may depend on the nature of the contribution. For example
   (a) If the contribution is primarily a new algorithm, the paper should make it clear how to reproduce that algorithm.
   (b) If the contribution is primarily a new model architecture, the paper should describe the architecture clearly and fully.
   (c) If the contribution is a new model (e.g., a large language model), then there should either be a way to access this model for reproducing the results or a way to reproduce the model (e.g., with an open-source dataset or instructions for how to construct the dataset).
   (d) We recognize that reproducibility may be tricky in some cases, in which case authors are welcome to describe the particular way they provide for reproducibility. In the case of closed-source models, it may be that access to the model is limited in some way (e.g., to registered users), but it should be possible for other researchers to have some path to reproducing or verifying the results.

5. **Open access to data and code**

   Question: Does the paper provide open access to the data and code, with sufficient instructions to faithfully reproduce the main experimental results, as described in supplemental material?

   Answer: [Yes]

Justification: Code is submitted with requirements.txt and a README.txt listing every command, and is available at https://github.com/froggomelo/nlp-project-2. The newspaper text is openly available from the archive (Schoots et al., 2023); the language-split files we used were provided with the course, and the modern corpora and encoders download automatically.

Guidelines:

- The answer NA means that paper does not include experiments requiring code.
- Please see the NeurIPS code and data submission guidelines (https://nips.cc/public/guides/CodeSubmissionPolicy) for more details.
- While we encourage the release of code and data, we understand that this might not be possible, so "No" is an acceptable answer. Papers cannot be rejected simply for not including code, unless this is central to the contribution (e.g., for a new open-source benchmark).
- The instructions should contain the exact command and environment needed to run to reproduce the results. See the NeurIPS code and data submission guidelines (https://nips.cc/public/guides/CodeSubmissionPolicy) for more details.
- The authors should provide instructions on data access and preparation, including how to access the raw data, preprocessed data, intermediate data, and generated data, etc.
- The authors should provide scripts to reproduce all experimental results for the new proposed method and baselines. If only a subset of experiments are reproducible, they should state which ones are omitted from the script and why.
- At submission time, to preserve anonymity, the authors should release anonymized versions (if applicable).
- Providing as much information as possible in supplemental material (appended to the paper) is recommended, but including URLs to data and code is permitted.

6. **Experimental Setting/Details**

Question: Does the paper specify all the training and test details (e.g., data splits, hyperparameters, how they were chosen, type of optimizer, etc.) necessary to understand the results?

Answer: [Yes]

Justification: Methods give the 357/77/77 train/validation/test issues stratified by year, and k and C chosen on validation. The code and README.txt give the rest: the kNN neighbour count chosen on validation, normaliser thresholds fixed in advance to favour precision, the E5 pair threshold of 0.857 (which keeps the same share of candidate pairs as 0.55 for LaBSE), and class-balanced logistic regression (scikit-learn, L-BFGS).

Guidelines:

- The answer NA means that the paper does not include experiments.
- The experimental setting should be presented in the core of the paper to a level of detail that is necessary to appreciate the results and make sense of them.
- The full details can be provided either with the code, in appendix, or as supplemental material.

7. **Experiment Statistical Significance**

Question: Does the paper report error bars suitably and correctly defined or other appropriate information about the statistical significance of the experiments?

Answer: [Yes]

Justification: The main test results carry 95% bootstrap confidence intervals over test issues (1,000 resamples), and Table S2 marks the rows that have them. The H2 equivalence test uses 90% intervals against a ± 0.02 margin, and synthetic results show the standard deviation over five seeds (Figure 1).

Guidelines:

- The answer NA means that the paper does not include experiments.
- The authors should answer "Yes" if the results are accompanied by error bars, confidence intervals, or statistical significance tests, at least for the experiments that support the main claims of the paper.
- The factors of variability that the error bars are capturing should be clearly stated (for example, train/test split, initialization, random drawing of some parameter, or overall run with given experimental conditions).

- The method for calculating the error bars should be explained (closed form formula, call to a library function, bootstrap, etc.)
- The assumptions made should be given (e.g., Normally distributed errors).
- It should be clear whether the error bar is the standard deviation or the standard error of the mean.
- It is OK to report 1-sigma error bars, but one should state it. The authors should preferably report a 2-sigma error bar than state that they have a 96% CI, if the hypothesis of Normality of errors is not verified.
- For asymmetric distributions, the authors should be careful not to show in tables or figures symmetric error bars that would yield results that are out of range (e.g. negative error rates).
- If error bars are reported in tables or plots, The authors should explain in the text how they were calculated and reference the corresponding figures or tables in the text.

8. **Experiments Compute Resources**

   Question: For each experiment, does the paper provide sufficient information on the computer resources (type of compute workers, memory, time of execution) needed to reproduce the experiments?

   Answer: [Yes]

   Justification: README.txt: embedding takes about one hour on a free Kaggle notebook with a Tesla T4 GPU (15 GB), and every other stage runs on a CPU, with the analysis taking about 15 minutes per encoder.

   Guidelines:

   - The answer NA means that the paper does not include experiments.
   - The paper should indicate the type of compute workers CPU or GPU, internal cluster, or cloud provider, including relevant memory and storage.
   - The paper should provide the amount of compute required for each of the individual experimental runs as well as estimate the total compute.
   - The paper should disclose whether the full research project required more compute than the experiments reported in the paper (e.g., preliminary or failed experiments that didn't make it into the paper).

9. **Code Of Ethics**

   Question: Does the research conducted in the paper conform, in every respect, with the NeurIPS Code of Ethics `https://neurips.cc/public/EthicsGuidelines`?

   Answer: [Yes]

   Justification: We use public-domain historical newspapers and no human subjects, and the Discussion covers how modernising rewrites the names people used in 1884, including God (Tixo).

   Guidelines:

   - The answer NA means that the authors have not reviewed the NeurIPS Code of Ethics.
   - If the authors answer No, they should explain the special circumstances that require a deviation from the Code of Ethics.
   - The authors should make sure to preserve anonymity (e.g., if there is a special consideration due to laws or regulations in their jurisdiction).

10. **Broader Impacts**

    Question: Does the paper discuss both potential positive societal impacts and negative societal impacts of the work performed?

    Answer: [Yes]

    Justification: On the positive side, the paper measures how a bilingual newspaper's sections differ and asks future archive studies to report the floor under both spellings. On the negative side, the Discussion notes that modernising rewrites how people named themselves, and that tools reading English more easily may tempt researchers to study only the English half.

    Guidelines:

    - The answer NA means that there is no societal impact of the work performed.

- If the authors answer NA or No, they should explain why their work has no societal impact or why the paper does not address societal impact.
- Examples of negative societal impacts include potential malicious or unintended uses (e.g., disinformation, generating fake profiles, surveillance), fairness considerations (e.g., deployment of technologies that could make decisions that unfairly impact specific groups), privacy considerations, and security considerations.
- The conference expects that many papers will be foundational research and not tied to particular applications, let alone deployments. However, if there is a direct path to any negative applications, the authors should point it out. For example, it is legitimate to point out that an improvement in the quality of generative models could be used to generate deepfakes for disinformation. On the other hand, it is not needed to point out that a generic algorithm for optimizing neural networks could enable people to train models that generate Deepfakes faster.
- The authors should consider possible harms that could arise when the technology is being used as intended and functioning correctly, harms that could arise when the technology is being used as intended but gives incorrect results, and harms following from (intentional or unintentional) misuse of the technology.
- If there are negative societal impacts, the authors could also discuss possible mitigation strategies (e.g., gated release of models, providing defenses in addition to attacks, mechanisms for monitoring misuse, mechanisms to monitor how a system learns from feedback over time, improving the efficiency and accessibility of ML).

11. **Safeguards**

    Question: Does the paper describe safeguards that have been put in place for responsible release of data or models that have a high risk for misuse (e.g., pretrained language models, image generators, or scraped datasets)?

    Answer: [NA]

    Justification: We release no models or scraped data with a high risk of misuse; outputs are statistics and excerpts of public-domain text.

    Guidelines:

    - The answer NA means that the paper poses no such risks.
    - Released models that have a high risk for misuse or dual-use should be released with necessary safeguards to allow for controlled use of the model, for example by requiring that users adhere to usage guidelines or restrictions to access the model or implementing safety filters.
    - Datasets that have been scraped from the Internet could pose safety risks. The authors should describe how they avoided releasing unsafe images.
    - We recognize that providing effective safeguards is challenging, and many papers do not require this, but we encourage authors to take this into account and make a best faith effort.

12. **Licenses for existing assets**

    Question: Are the creators or original owners of assets (e.g., code, data, models), used in the paper, properly credited and are the license and terms of use explicitly mentioned and properly respected?

    Answer: [Yes]

    Justification: The Imvo texts are in the public domain (Schoots et al., 2023). Licences from the Hugging Face cards: Vukuzenzele CC BY 4.0, isiXhosa Wikipedia CC BY-SA 3.0 and GFDL, LaBSE Apache 2.0, multilingual-E5 and the XLM-R tokeniser MIT. OPUS-100's card lists its licence as unknown, so we use it only to count word frequencies and redistribute none of its text. Dataset revisions are recorded in results/lexicon/lexicon_sources.json.

    Guidelines:

    - The answer NA means that the paper does not use existing assets.
    - The authors should cite the original paper that produced the code package or dataset.
    - The authors should state which version of the asset is used and, if possible, include a URL.
    - The name of the license (e.g., CC-BY 4.0) should be included for each asset.
    - For scraped data from a particular source (e.g., website), the copyright and terms of service of that source should be provided.

- If assets are released, the license, copyright information, and terms of use in the package should be provided. For popular datasets, `paperswithcode.com/datasets` has curated licenses for some datasets. Their licensing guide can help determine the license of a dataset.
- For existing datasets that are re-packaged, both the original license and the license of the derived asset (if it has changed) should be provided.
- If this information is not available online, the authors are encouraged to reach out to the asset's creators.

13. **New Assets**

    Question: Are new assets introduced in the paper well documented and is the documentation provided alongside the assets?

    Answer: [Yes]

    Justification: The code and normaliser are documented in README.txt, and the hand labels are in data/pair_labels.csv.

    Guidelines:

    - The answer NA means that the paper does not release new assets.
    - Researchers should communicate the details of the dataset/code/model as part of their submissions via structured templates. This includes details about training, license, limitations, etc.
    - The paper should discuss whether and how consent was obtained from people whose asset is used.
    - At submission time, remember to anonymize your assets (if applicable). You can either create an anonymized URL or include an anonymized zip file.

14. **Crowdsourcing and Research with Human Subjects**

    Question: For crowdsourcing experiments and research with human subjects, does the paper include the full text of instructions given to participants and screenshots, if applicable, as well as details about compensation (if any)?

    Answer: [NA]

    Justification: No crowdsourcing or research with human subjects. A fluent isiXhosa speaker labelled 24 sampled translated pairs herself and reviewed 27 more that a team member had labelled with AI translations and suggested labels (translation, partial or different), following a written labelling guide. No personal data was collected.

    Guidelines:

    - The answer NA means that the paper does not involve crowdsourcing nor research with human subjects.
    - Including this information in the supplemental material is fine, but if the main contribution of the paper involves human subjects, then as much detail as possible should be included in the main paper.
    - According to the NeurIPS Code of Ethics, workers involved in data collection, curation, or other labor should be paid at least the minimum wage in the country of the data collector.

15. **Institutional Review Board (IRB) Approvals or Equivalent for Research with Human Subjects**

    Question: Does the paper describe potential risks incurred by study participants, whether such risks were disclosed to the subjects, and whether Institutional Review Board (IRB) approvals (or an equivalent approval/review based on the requirements of your country or institution) were obtained?

    Answer: [NA]

    Justification: No human subjects: the texts are newspapers from 1884–1894, and the archive's authors note that no ethics review was needed because the materials discuss no living subjects (Schoots et al., 2023).

    Guidelines:

    - The answer NA means that the paper does not involve crowdsourcing nor research with human subjects.
    - Depending on the country in which research is conducted, IRB approval (or equivalent) may be required for any human subjects research. If you obtained IRB approval, you should clearly state this in the paper.
    - We recognize that the procedures for this may vary significantly between institutions and locations, and we expect authors to adhere to the NeurIPS Code of Ethics and the guidelines for their institution.
    - For initial submissions, do not include any information that would break anonymity (if applicable), such as the institution conducting the review.

Wits University Faculty of Science post-graduate student AI declaration

I understand that the use of generative AI tools (such as ChatGPT or similar) without explicitly declaring such use constitutes a form of plagiarism and is classified by Wits University as academic misconduct.

I declare that in the course of conducting the research towards my degree or in the preparation of this thesis/dissertation/research report (select one by marking with an X):

I **did not** make use of generative AI tools ☐

I **did** make use of generative AI tools for the following (tick all that apply):

1.	Idea Generation (research problem/design, hypothesis)	X
2.	Sourcing Related Work (summarising, identifying sources)	X
3.	Methods and Experiment Design (experiment setup, model tuning)	X
4.	Data Analysis (presentation, coding, interpretation)	X
5.	Theoretical Development (theorem proving, conceptual analysis)	
6.	Code Development (generating algorithms, writing scripts)	X
7.	Presentation (rendering graphics, formatting)	X
8.	Editing (grammar, readability)	X
9.	Writing (text generation, document structuring)	X
10.	Citation Formatting (structuring, organising)	X

If other uses were involved, please specify below:

Generative AI tool used (list all)	**Used for?**

If generative AI tools were used as an integral part of the experimental design or in the direct execution of my research, I confirm that details of this use are clearly outlined in the relevant experimental/methodology chapters of my thesis/dissertation/research report.

Student number:

2545080 2680747 2588206

Candidate signature:





Date: 30/10/2026

Supplementary material

Two tables and a figure that support claims in the body, from the saved results in `results/`.

Table S1: Spelling variants on test issues with LaBSE. The last column is each variant's change in Δ from original (raw) spelling, with its 90% CI (1,000 resamples of test issues).

Variant	All	Floor	Δ	$\Delta - \Delta_{\text{raw}}$ (90% CI)
raw	.786	.590	.196	–
modern	.784	.571	.214	$+$.018 ($-$.003, $+$.037)
aspiration only	.783	.569	.213	$+$.017 ($-$.001, $+$.038)
prefix only	.786	.600	.187	$-$.009 ($-$.024, $+$.006)

Table S2: Δ across checks, on test issues with LaBSE and original (raw) spelling unless stated. Brackets give the 95% CI where available (1,000 resamples of test issues). The top five rows use the same classifier and differ only in which pairs give the floor, and the two in bold set the range of Δ in Results. Pairs matched on shared names and numbers are found without embeddings. The topic-only rule reruns the clustering in Figure S1, so cluster sizes differ slightly. kNN is a cosine k-nearest-neighbour classifier.

Setting	All	Floor	Δ (95% CI)
Translated pairs hand-checked as full translations (16 pairs)	**.786**	**.500**	**.286 (.142, .412)**
Pairs matched on shared names and numbers, but chunks not in translated pairs (252 chunks)	**.786**	**.670**	**.117 (.066, .169)**
All translated pairs, similarity $\geq$0.55 (311 pairs)	.786	.590	.196 (.167, .224)
Pairs matched on shared names and numbers (217 pairs)	.786	.627	.159 (.128, .191)
Translated pairs with similarity $\geq$0.60 (161 pairs)	.786	.547	.239 (.203, .276)
Topic-only rule	.661	.523	.138
E5 encoder	.877	.618	.258 (.216, .304)
kNN classifier, E5 encoder	.599	.553	.046

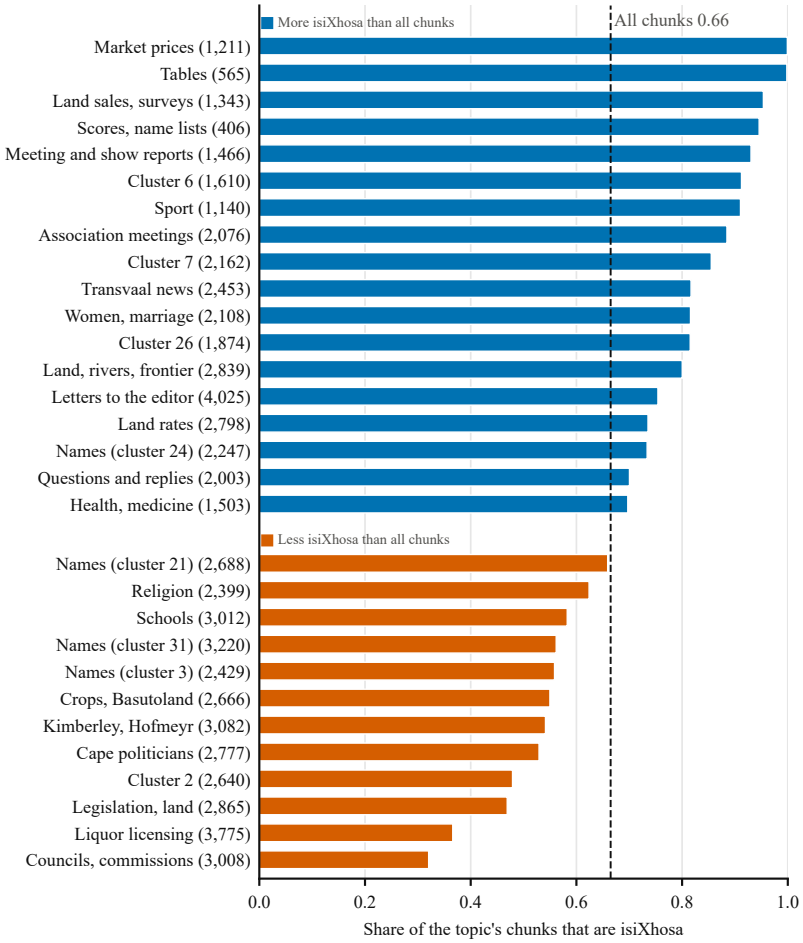


Figure S1: What each section usually wrote about. Each bar is one of 30 topics, with its chunk count, and shows the share of its chunks that are isiXhosa. Blue topics are above the overall share (0.66) and orange ones are below it. Topics are k-means clusters of the language-removed embeddings, named by the authors from their most distinctive words. Clusters labelled *names* hold mostly personal and place names, and numbered ones have no clear subject. The topics overlap (silhouette score 0.008, where 0 means fully overlapping), so they describe content rather than measure it.